

AnyBlox: Decoupling Storage Formats from Data Processing Systems

Liam Sanft

liam.sanft@mailbox.tu-dresden.de

TU Dresden

Germany, Dresden

Abstract

Open storage formats in data lakes have separated data from the systems that process it, reducing vendor lock-in. However, reading a format still requires a decoder for that format to be present in every consuming system. For N systems and M formats, this results in $N \times M$ independent decoder implementations, which raise maintenance costs, increase the risk of security vulnerabilities, and discourage the adoption of new encodings. To address this, Gieniec et al. (2025) propose AnyBlox, a framework that moves the decoder into the data itself: each dataset ships with its own decoder as sandboxed WebAssembly bytecode. With this approach, any engine that implements the framework once can read a format without prior knowledge, reducing the integration burden to $N + M$. The sandbox overhead remains negligible in end-to-end workloads. Furthermore, more efficient encodings outperform established formats in both scan latency and compression, directly addressing the maintenance burden and format ossification caused by the $N \times M$ problem.

1 Introduction

For decades, database systems functioned as tightly coupled units, with a single system controlling storage format, processing logic, and data access. External systems submitted queries and received results; internal operations remained hidden. With complete control over the physical layout, systems can optimise data storage and processing within their own boundaries. This model was sufficient as long as data did not cross system boundaries [19].

However, the volume and variety of collected data have grown drastically. Data now arrives continuously from a growing number of sources, in formats suited to their respective domains rather than to any particular database system. As a result, the database landscape has shifted from a single system controlling all aspects of data management to a heterogeneous, modular ecosystem of specialised systems, each handling a distinct part of the data pipeline using a shared storage layer. Open file formats such as Apache Parquet¹ and Apache ORC², stored in data lakes [3], are widely adopted in this context. Beyond these, data scientists increasingly analyse data using domain-specific encodings from fields such as high-energy physics and bioinformatics [8].

Independently, research in data compression and encoding continues to produce formats that surpass established standards for storage efficiency and query throughput. Yet, these advances rarely appear in production systems, because each system must independently implement support for every new format it wants to read [8].




Figure 1: (Top) A monolithic DBMS controls its query engine, storage engine, and internal format; clients interact only through queries. (Bottom) In a data lake, multiple query engines access a shared pool of open file formats. Since each must implement a decoder for each format, N engines and M formats require $N \times M$ independent implementations.

Both trends lead to the same structural problem. Codd’s principle of physical data independence [6] assumes a system controls its storage format and shields applications from physical details. In a shared storage layer, this does not apply. The data producer determines the format, so any system that reads a format must implement its own decoder. With N systems and M formats, there are $N \times M$ independent implementations to develop and maintain.

This results in higher implementation costs, greater maintenance effort, and increased security risks, as each implementation may contain bugs. For example, a high-severity vulnerability in PostgreSQL’s³ JSON support went undetected for over five years [17]. This initial hurdle discourages systems from supporting new formats altogether.

The consequence is a self-sustaining cycle: systems do not adopt a new file format because its user base is too small, and the user base remains small because the format lacks support. Gieniec et al. (2025) identify this as a core structural problem because an encoding must be popular enough to justify support, but cannot gain popularity without broad support.

¹Apache Parquet, <https://parquet.apache.org> (accessed: May 22, 2026)

²Apache ORC, <https://orc.apache.org/> (accessed: May 22, 2026)

³PostgreSQL, <https://www.postgresql.org/> (accessed: May 22, 2026)

This dynamic leads to format ossification, where datasets use outdated formats just to stay compatible, even when better encodings exist [12]. Changing formats is difficult because it risks breaking compatibility with many systems that use older specifications.

But what could solve this combinatorial integration burden? An abstraction layer between storage formats and processing systems reduces the $N \times M$ problem to $N + M$. Each system implements the interface once, and each format provides a single implementation against it. Gienieccko et al. (2025) propose exactly this and present AnyBlox as their realisation.

2 Self-Decoding Data: The AnyBlox Design

A future-proof data access solution must satisfy two conceptual requirements simultaneously. First, it must introduce an abstraction layer that decouples format internals from processing systems and system internals from decoders, thereby reducing the $N \times M$ integration burden to $N + M$. Second, this abstraction must not preclude direct file access, since routing data through an intermediary process introduces transfer overhead that can dominate total runtime. These two goals are in tension: an abstraction layer must be located somewhere, and every candidate location either breaks direct access or reintroduces format-specific coupling [8].

Existing approaches each resolve this tension in favor of one side. Native integration preserves direct file access and achieves the best performance through tight coupling with host internals, but requires a separate decoder implementation per format per system, thus reproducing the $N \times M$ problem. Extension mechanisms reduce the per-format implementation effort, but each host exposes its own extension API, so a decoder written for system A cannot be reused in system B without a near-complete rewrite. Isolated extensions, such as containerized user-defined functions [18] or remote cloud user-defined code execution solutions provide both isolation and portability, but introduce a data transfer bottleneck between the host and the isolated decoder process [8].

Given these constraints, Gienieccko et al. (2025) argue that there is only one location for the decoder that satisfies both requirements: inside the data file itself. Packaging the decoder with the data eliminates the transfer bottleneck, since decoder and data reside in the same address space, while the abstraction is preserved because the host need not implement any format-specific logic. This is the core idea behind self-decoding data and the design premise of AnyBlox.

2.1 WebAssembly as the Decoder Runtime

A self-decoding dataset requires a decoder runtime that satisfies four criteria: **Portability**, **Security**, **Performance** and **Extensibility**.

Gienieccko et al. (2025) identify WebAssembly (Wasm) as a runtime that satisfies all four simultaneously. Wasm is a portable low-level bytecode and virtual machine specification, designed as an abstraction over modern hardware rather than a commitment to any specific programming model, making it language-, hardware-, and platform-independent [9].

2.1.1 Portability. A decoder implemented in Wasm runs on a wide variety of host architectures (e.g. browsers, mobile devices, embedded, and server workstations) [8]. For AnyBlox, this means a format

author writes and compiles a decoder once, and any host can execute it via `libanyblox`, independent of the underlying hardware or operating system.

2.1.2 Security. The WebAssembly specification provides machine-verifiable memory safety and isolation guarantees [21]. AnyBlox reinforces this by being implemented predominantly in safe Rust, guaranteeing memory and thread safety. All unsafe code is confined to a single module handling Wasm memory maps, keeping it easily auditable [8].

2.1.3 Performance. Wasm is JIT-compiled to the host’s native instruction set at runtime, making near-native performance attainable. Additionally, AnyBlox avoids unnecessary data copying through a custom memory management scheme. Gienieccko et al. (2025) note that AnyBlox can passively benefit from future improvements in Wasm compiler technology, since all JIT and sandboxing details are encapsulated in `libanyblox` and not exposed in the public API [8].

2.1.4 Extensibility. Since most general-purpose programming languages provide a Wasm toolchain, format authors can implement decoders in their language of choice without restrictions. The main practical obstacle is SIMD instructions: converting platform-specific SIMD intrinsics to Wasm’s instruction set requires manual developer effort. However, since Wasm defines SIMD at the bytecode level, the resulting modules remain fully portable across architectures [8].

2.2 Reading a Dataset: From File to Query Pipeline

When a Host wants to read an AnyBlox dataset, it first opens the `.any` file and reads its metadata during query planning. The metadata provides the Host with the schema, the total number of rows, and a recommended batch size, all of which are needed to partition the workload across threads and schedule the scan before decoding begins.

The Host then hands the metadata and `.any` file to `libanyblox`, which extracts the Wasm Decoder module, JIT-compiles it, and caches the result. Subsequent uses of the same Decoder therefore incur no compilation cost. Following this, AnyBlox sets up the execution environment for the Decoder and returns a job handle to the Host.

This procedure describes standalone AnyBlox, where the dataset resides in its own `.any` file and the metadata is limited to schema, row count, and batch size. AnyBlox can alternatively be embedded into an existing host format such as Parquet, in which case the Host is the format’s own decoding library and the embedded path inherits that format’s statistics and row-group filtering. The decoding pipeline described below is identical in both cases. The two modes differ only in the metadata available to the Host, a distinction that becomes relevant for predicate pushdown (Section 3.3).

From this point, the Host repeatedly requests arbitrary tuple ranges from that handle. Modern query engines process data in batches to amortize operator communication overhead, so rather than decoding the entire dataset at once, the Host issues successive calls for ranges of tuples [11]. AnyBlox delegates each such request into the sandboxed Wasm module and returns the decoded

result to the Host. The sole function through which this happens is `decode_batch` [8]:

- `i32 data`: pointer to the start of the encoded data in Wasm linear memory
- `i32 data_length`: length of the encoded data in bytes
- `i32 start_tuple`: ID of the first tuple to decode
- `i32 tuple_count`: number of tuples to decode
- `i32 state`: pointer to the state page in Wasm linear memory
- `i64 projection_mask`: bitmask specifying which columns to decode, allowing the Host to skip columns that are not needed for the current query

The data parameter points to the encoded dataset inside the Decoder’s execution environment, but how does the dataset get there in the first place? A naive solution would be to copy the entire dataset into a memory region accessible to the Decoder, but this would incur a heavy initialization cost proportional to the dataset size. Instead, AnyBlox maps the dataset file directly into the Decoder’s linear memory via `mmap`, a system call that maps a file directly into the virtual address space of a process without copying it, so physical memory is only allocated when a page is first accessed.

Wasm’s linear memory is a contiguous, byte-addressable memory region starting at address `0x0` that represents the entire address space accessible to the Decoder, isolated from the rest of the host process. AnyBlox controls its layout: the encoded dataset is mapped into it via `mmap`, and the state page is placed immediately after. The data parameter is therefore simply a pointer to the start of the mapped dataset within this region. AnyBlox reuses the linear memory region across successive jobs on the same thread, avoiding repeated initialization costs.

The state parameter points to a page within the linear memory that is zero-initialized at the start of every job. A Decoder can use it to cache metadata or memory allocations across successive `decode_batch` calls. For example, a run-end encoding Decoder that performs a binary search on the first call can store its current position on the state page and resume directly on the next call, skipping the search entirely. Because the state page resets at the start of each job, Decoders are idempotent: identical parameters on the same dataset always produce equivalent results.

Once decoding is complete, `decode_batch` returns a pointer to an Apache Arrow Record Batch within the linear memory. Arrow⁴ is a columnar in-memory format developed by the Apache Software Foundation: rather than storing data row by row, it organizes each column as a contiguous array in memory, which allows query engines to process large amounts of data efficiently without unnecessary memory accesses. A Record Batch is a collection of such arrays, all of the same length, representing a horizontal slice of a table. Arrow standardizes the representation of common SQL types (e.g., integers, floating-point numbers, strings, dates, and more) and is supported by a wide range of query engines. The choice of Arrow as the uniform output format of AnyBlox follows the same principle as other intermediate representations in systems design, because by agreeing on a common language between two components, the number of required translations reduces from $N \times M$ to $N + M$. Every

⁴Apache Arrow, <https://arrow.apache.org/> (accessed: June 2, 2026)




Figure 2: Total query latency (median, summed over 16 TPC-H queries at scale factor 20) in DataFusion, split into scan time (dark) and operator time (light). Percentages denote the scan share of total latency. Data source: [8].

Decoder, regardless of the underlying encoding scheme, produces an Arrow Record Batch, and every Host that speaks Arrow can consume it without any AnyBlox-specific logic. AnyBlox translates the Wasm-internal pointer into a native address and passes the batch to the Host, where it enters the query pipeline. The degree of integration varies by system: DataFusion [13] uses Arrow as its primary internal representation throughout, making the transfer zero-copy end-to-end. DuckDB [16] and Spark⁵ do not use Arrow internally, but both have built-in support for reading Arrow batches and converting them into their respective internal representations on the fly. This conversion logic is part of the mainline code of both systems and requires no AnyBlox-specific adaptation [8].

3 Discussion

The preceding sections introduced the AnyBlox architecture and its decoding model. Whether this design is practical depends on three questions:

- (1) What overhead does the Wasm sandbox introduce?
- (2) How much effort does integrating the framework into an engine require?
- (3) Where does the approach reach its limits?

Gienieczo et al. (2025) report benchmarks and integration experience that cover both favorable and unfavorable cases for AnyBlox.

3.1 Performance Evaluation

Sandboxing a decoder inside Wasm adds a layer of indirection between the data and the query engine. Gienieczo et al. (2025) measure the resulting overhead directly. They run TPC-H at scale factor 20 in DataFusion [13] and vary the representation of `lineitem`, the largest table in the benchmark and therefore the one whose scan dominates the runtime, across three configurations:

- Parquet** read by DataFusion’s native reader with Snappy⁶ compression;
- AnyBlox Vortex** a Vortex file decoded inside the AnyBlox sandbox;
- Native Vortex** the same Vortex decoder compiled natively without any sandbox.

⁵Apache Spark, <https://spark.apache.org/> (accessed: June 2, 2026)

⁶Snappy, <https://github.com/google/snappy/> (accessed: June 15, 2026)

The first variant isolates the effect of the format; the latter two isolate the effect of the sandbox alone. Predicate pushdown was disabled throughout, so the numbers reflect scan and operator cost only. Vortex⁷ is a recent columnar format.

The benchmark yields two results, summarized in Figure 2. Both Vortex variants achieve over 2× lower scan latency than Parquet while also reaching a better compression rate. For the Parquet reader, decoding accounts for nearly half of the processing time, while Vortex spends around 80% of its processing time in relational operators. The overhead of the sandbox itself is small. Sandboxed Vortex decoding adds 5% to the scan latency over native decoding, which amounts to a 1% increase in overall query latency and which the authors report as within measurement noise. Across this end-to-end workload, the cost of the sandbox is therefore negligible, and AnyBlox provides DataFusion with an encoding that outperforms its own native Parquet reader [8].

The microbenchmarks show where the overhead of the sandbox becomes visible. FSST [5] is a lightweight string compression scheme. Gienieczko et al. (2025) benchmark the decoding of three FSST-compressed string columns from the Taxpayer dataset of the Public BI Benchmark [7]. The 1.7 GB AnyBlox file expands into 2.5 GB of Arrow batches in memory. In this isolated scenario, the Wasm Decoder is over 43% slower than the same decoder compiled natively. The cause is architectural rather than incidental: the FSST decoder relies heavily on 8-byte reads, but Wasm uses a 32-bit address space by default, so each 8-byte read is emulated with two 4-byte reads. The authors describe this as a fundamental limitation of the current design. The 64-bit Wasm proposal would remove it, but it is incompatible with the software fault isolation mechanism on which AnyBlox’s security guarantees rest, since that mechanism assumes a Decoder can address at most 8 GiB of memory [8].

Taken together, the benchmarks place the cost of the sandbox in perspective, though the evidence differs in strength across the two cases. For a complex decoder such as Vortex, the overhead is measured directly: in the end-to-end DataFusion workload it stays within measurement noise, and the more efficient encoding lets AnyBlox outperform the native Parquet reader. For the FSST worst case, the slowdown appears clearly in the microbenchmark, but its effect on full queries is not measured directly. Since the scan accounts for only about a fifth of processing time in the DataFusion workload, a decoding slowdown can affect only that share of the runtime. The microbenchmark therefore isolates the decoder under the conditions least favorable to it.

3.2 Integration Effort

The second claim AnyBlox makes is portability, meaning an engine can support the framework without a large implementation effort. Gienieczko et al. (2025) integrated AnyBlox into four systems spanning different execution paradigms (Umbra [15], DuckDB, DataFusion and Spark), all done by a single programmer with no prior experience in any of the host codebases. Table 1 lists the resulting effort. No integration required more than 1305 lines of code, and each one only added a new scan operator while leaving the rest of the engine, namely its optimizer, operators and execution layer, unchanged [8].

⁷Vortex, <https://vortex.dev/> (accessed: June 15, 2026)

Table 1: Effort of integrating AnyBlox into each host system. Data source: [8].

System	Lines of code	Person-days
DataFusion	461	2
DuckDB	586	3
Spark	756	15
Umbra	1305	10

The authors caution that lines of code and integration time are not rigorous measures of software complexity, so these figures indicate rather than prove that the framework ports easily. With that caveat, they are the empirical counterpart to the $N \times M$ argument: a one-time integration lets an engine read any format providing a decoder.

3.3 Limitations

Two of AnyBlox’s limitations follow from the design itself rather than from the current implementation. The first concerns predicate pushdown. Predicate pushdown evaluates filter conditions during the scan phase, so that rows failing a condition are discarded before they are decoded and passed to operators such as aggregation or joins. Decoding is computationally expensive, so prior filtering avoids processing data that would be discarded anyway [19].

Standalone AnyBlox provides no predicate pushdown API for row filtering. It supports projection pushdown, which restricts a scan to the requested columns, but not the filtering of individual rows. The reason is a deliberate design choice: filtering individual rows in a format-agnostic framework requires an equally format-agnostic and system-independent predicate expression format, which AnyBlox does not define. This reflects the same design philosophy visible elsewhere in the framework, that statistics and indexes should be decoupled from the storage format and managed by the query engine or a dedicated acceleration layer. The authors point to Substrait⁸ as ongoing industry work toward a standardized predicate expression format, which could close this gap in a future revision [8].

This limitation is confined to standalone use. When an AnyBlox encoding is embedded into a host format such as Parquet at the row-group or column level, the host’s existing row-group and column filtering continues to apply, so the embedded path retains the pushdown that the host already performs [8].

The second limitation concerns memory. Standalone AnyBlox maps the entire dataset into the Wasm linear memory region via `mmap` before decoding begins. The current implementation uses `wasm32`, which provides a 32-bit address space. Each mapped file is therefore limited to 4 GB. The authors note that file sizes in typical analytical deployments are commonly kept well below 4 GiB to facilitate cheap updates, so the restriction may have limited impact in practice. A migration to `wasm64` would remove the limit, since Memory64 has been part of the WebAssembly standard since 2025⁹. However, the performance impact of this migration in the

⁸Substrait, <https://substrait.io/> (accessed: June 15, 2026)

⁹WebAssembly Memory64, <https://github.com/WebAssembly/memory64> (accessed: June 15, 2026)

context of AnyBlox is not yet studied, and wasm64 is incompatible with the software fault isolation mechanism on which AnyBlox’s security guarantees rest. The mmap requirement also means that column projections cannot be pushed to remote files, because the entire dataset must reside in memory before the decoder is invoked, which is a relevant constraint in data lake deployments where data typically resides on object stores. Partial remote reads via page-fault handling are sketched as future work but not yet implemented [8].

4 Conclusion

Reading an unfamiliar storage format normally requires a reader written for that format inside every system that wants to consume it. AnyBlox removes this requirement by moving the reader into the data. A decoder travels with the dataset as sandboxed WebAssembly bytecode, and any engine that implements a single scan interface can decode the format without knowing it in advance. The cost of this design is the overhead of the sandbox, which the measurements of Gienieccko et al. (2025) show to be negligible for the end-to-end workloads tested.

The authors frame the underlying problem as an $N \times M$ problem: N systems each supporting M formats incur $N \times M$ implementation and maintenance effort [8]. An abstraction layer between the two sides reduces this to $N + M$, since each engine implements the interface once and each format provides one decoder. The authors draw this analogy explicitly from compiler construction, citing the language-independent intermediate representation of LLVM [14], which decouples M language frontends from N target backends through a shared representation. AnyBlox applies the same decoupling between storage formats and query engines. The comparison is an analogy rather than evidence: it motivates the design but does not establish that the result transfers from compilers to data formats.

This decoupling is not an isolated idea but part of a broader move toward composable database systems, away from large, monolithic designs. The authors place AnyBlox alongside several efforts that each modularise a different layer, among them a query engine built on composable design principles [13], execution-plan frameworks made extensible through sub-operators [4, 10], portable query optimizers [1, 2], and system-independent persistence mechanisms [20]. AnyBlox contributes the storage-format layer to this picture, and one of these efforts connects directly to it: the predicate pushdown limitation discussed in Section 3.3, where standalone AnyBlox deliberately defines no row-filtering predicate format, is a gap that Substrait, an emerging cross-engine standard for query plans, could close. The components are complementary, each standardising one interface between otherwise independent systems [8].

As a precedent that such modular frameworks can succeed, the authors invoke LLVM, which reshaped the design of compilers by establishing a shared, open representation [8, 14]. AnyBlox shows that the same principle could work for storage formats. Its evaluation demonstrates that a format-agnostic decoding layer can be built with low overhead and integrated into engines across different execution paradigms with modest effort. The LLVM comparison remains an analogy, and the evidence so far comes from a single evaluation by the original authors, so whether modular stacks come to displace integrated systems is a question that broader adoption

and independent study will decide. What the paper establishes is that the decoding layer is no longer the obstacle.

References

- [1] Rana Alotaibi, Yuanyuan Tian, Stefan Grafberger, Jesús Camacho-Rodríguez, Nicolas Bruno, Brian Kroth, Sergiy Matusevych, Ashvin Agrawal, Mahesh Behera, Ashit Gosalia, Cesar Galindo-Legaria, Milind Joshi, Milan Potocnik, Beysim Sezgin, Xiaoyu Li, and Carlo Curino. 2024. Towards Query Optimizer as a Service (QOaaS) in a Unified LakeHouse Ecosystem: Can One QO Rule Them All? <https://doi.org/10.48550/arXiv.2411.13704> arXiv:2411.13704 [cs.DB]
- [2] Christoph Anneser, Nesime Tatbul, David Cohen, Zhenggang Xu, Prithviraj Pandian, Nikolay Laptev, and Ryan Marcus. 2023. AutoSteer: Learned Query Optimization for Any SQL Database. *Proceedings of the VLDB Endowment* 16, 12 (Aug. 2023), 3515–3527. <https://doi.org/10.14778/3611540.3611544>
- [3] Michael Armbrust, Ali Ghodsi, Reynold Xin, and Matei Zaharia. 2021. Lakehouse: A New Generation of Open Platforms That Unify Data Warehousing and Advanced Analytics. (2021).
- [4] Maximilian Bandle and Jana Giceva. 2021. Database Technology for the Masses: Sub-Operators as First-Class Entities. *Proceedings of the VLDB Endowment* 14, 11 (July 2021), 2483–2490. <https://doi.org/10.14778/3476249.3476296>
- [5] Peter Boncz, Thomas Neumann, and Viktor Leis. 2020. FSST: Fast Random Access String Compression. *Proceedings of the VLDB Endowment* 13, 12 (Aug. 2020), 2649–2661. <https://doi.org/10.14778/3407790.3407851>
- [6] E. F. Codd. 1970. A Relational Model of Data for Large Shared Data Banks. *Commun. ACM* 13, 6 (June 1970), 377–387. <https://doi.org/10.1145/362384.362685>
- [7] Bogdan Ghit, Diego Tomé, and Peter Boncz. [n. d.]. White-Box Compression: Learning and Exploiting Compact Table Representations. ([n. d.]).
- [8] Mateusz Gienieccko, Maximilian Kuschewski, Thomas Neumann, Viktor Leis, and Jana Giceva. 2025. AnyBlox: A Framework for Self-Decoding Datasets. *Proceedings of the VLDB Endowment* 18, 11 (July 2025), 4017–4031. <https://doi.org/10.14778/3749646.3749672>
- [9] Andreas Haas, Andreas Rossberg, Derek L. Schuff, Ben L. Titzer, Michael Holman, Dan Gohman, Luke Wagner, Alon Zakai, and JF Bastien. 2017. Bringing the Web up to Speed with WebAssembly. *SIGPLAN Not.* 52, 6 (June 2017), 185–200. <https://doi.org/10.1145/3140587.3062363>
- [10] Michael Jungmair and Jana Giceva. 2023. Declarative Sub-Operators for Universal Data Processing. *Proceedings of the VLDB Endowment* 16, 11 (July 2023), 3461–3474. <https://doi.org/10.14778/3611479.3611539>
- [11] Timo Kersten, Viktor Leis, Alfons Kemper, Thomas Neumann, Andrew Pavlo, and Peter Boncz. 2018. Everything You Always Wanted to Know about Compiled and Vectorized Queries but Were Afraid to Ask. *Proceedings of the VLDB Endowment* 11, 13 (Sept. 2018), 2209–2222. <https://doi.org/10.14778/3275366.3275370>
- [12] Laurens Kuiper. 2025. Query Engines: Gatekeepers of the Parquet File Format. <https://duckdb.org/2025/01/22/parquet-encodings.html>.
- [13] Andrew Lamb, Yijie Shen, Daniel Heres, Jayjeet Chakraborty, Mehmet Ozan Kabak, Liang-Chi Hsieh, and Chao Sun. 2024. Apache Arrow DataFusion: A Fast, Embeddable, Modular Analytic Query Engine. In *Companion of the 2024 International Conference on Management of Data (SIGMOD ’24)*. Association for Computing Machinery, New York, NY, USA, 5–17. <https://doi.org/10.1145/3626246.3653368>
- [14] C. Lattner and V. Adve. 2004. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *International Symposium on Code Generation and Optimization, 2004. CGO 2004*. IEEE, San Jose, CA, USA, 75–86. <https://doi.org/10.1109/CGO.2004.1281665>
- [15] Thomas Neumann and Michael Freitag. [n. d.]. Umbra: A Disk-Based System with In-Memory Performance. ([n. d.]).
- [16] Mark Raasveldt and Hannes Mühleisen. 2019. DuckDB: An Embeddable Analytical Database. In *Proceedings of the 2019 International Conference on Management of Data (SIGMOD ’19)*. Association for Computing Machinery, New York, NY, USA, 1981–1984. <https://doi.org/10.1145/3299869.3320212>
- [17] Red Hat, Inc. 2017. CVE-2017-15098. <https://nvd.nist.gov/vuln/detail/CVE-2017-15098>.
- [18] Karla Saur, Tara Mirmira, Konstantinos Karanasos, and Jesús Camacho-Rodríguez. 2022. Containerized Execution of UDFs: An Experimental Evaluation. *Proceedings of the VLDB Endowment* 15, 11 (July 2022), 3158–3171. <https://doi.org/10.14778/3551793.3551860>
- [19] Abraham Silberschatz, Henry F. Korth, and S. Sudarshan. 2020. *Database System Concepts* (seventh edition ed.). McGraw-Hill, New York, NY.
- [20] Emil Tsalapatis, Ryan Hancock, Rakeeb Hossain, and Ali José Mashtizadeh. 2024. MemSnap µCheckpoints: A Data Single Level Store for Fearless Persistence. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3. ACM, La Jolla CA USA*, 622–638. <https://doi.org/10.1145/3620666.3651334>
- [21] Conrad Watt. 2018. Mechanising and Verifying the WebAssembly Specification. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2018)*. Association for Computing Machinery, New York, NY, USA, 53–65. <https://doi.org/10.1145/3167082>