

AnyBlox: Decoupling Storage Formats from Data Processing Systems

Liam Sanft

liam.sanft@mailbox.tu-dresden.de

TU Dresden

Germany, Dresden

Abstract

Here comes the abstract.

1 Introduction

For decades, database systems operated as tightly coupled units in which storage format, processing logic, and data access were controlled by a single system. External systems submitted queries and received results, while everything in between remained internal concerns. With full control over the physical layout, systems could optimize both storage and processing of data freely within their own boundaries. As long as data resided within a single system, this model was sufficient [11].

However, the volume and variety of data collected has grown drastically. Data arrives continuously from a growing number of sources, in formats tailored to their respective domains rather than to any particular database system. As a result, the database landscape has evolved from a single system controlling all aspects of data management toward a heterogeneous, modular ecosystem of specialized systems, each handling a distinct part of the data pipeline and accessing a shared storage layer. Open file formats such as Apache Parquet¹ and Apache ORC², stored in data lakes [1], have gained wide adoption in this context. Beyond these, data scientists increasingly need to analyze data in domain-specific encodings from fields such as high-energy physics and bioinformatics [3].

Independently, research in data compression and encoding continues to produce formats that outperform established standards in storage efficiency and query throughput. These advances rarely reach production systems, however, because each system must independently implement support for every format it wishes to read [3].

Both trends converge on the same structural problem. Codd’s principle of physical data independence [2], which assumes that a system controls its own storage format and shields applications from its physical details, no longer holds in a shared storage layer. In a shared storage layer, this assumption no longer holds. The format is determined by whoever produced the data. Every system that wants to read a given format must implement its own decoder. With N systems and M formats in active use, this results in $N \times M$ independent implementations to develop and maintain [3].

This results in higher implementation cost and maintenance effort, and simultaneously increases the risk of security vulnerabilities, as each implementation may contain bugs, as illustrated by a high-severity vulnerability in PostgreSQL’s³ JSON support




Figure 1: (Top) A monolithic DBMS controls its query engine, storage engine, and internal format; clients interact only through queries. (Bottom) In a data lake, multiple query engines access a shared pool of open file formats. Since each engine must implement a decoder for each format, N engines and M formats require $N \times M$ independent implementations.

that went undetected for over five years [9]. This initial hurdle discourages systems from supporting new formats altogether.

The consequence is a self-reinforcing cycle: a new format is not adopted by systems because its user base is too small to justify the implementation effort, yet its user base remains small precisely because it is not widely supported. Gienieccko et al. (2025) identify this as a core structural problem: an encoding must be popular enough to justify the cost of support, but will not gain popularity without being widely supported.

This dynamic leads to format ossification: datasets continue to be produced using outdated format specifications solely to maintain compatibility with existing readers, even when more efficient encodings are available [6]. Evolution of a format is practically prevented, because any change risks breaking compatibility with the large number of systems that have independently implemented the old specification [3].

But what would be the solution for this combinatorial integration burden? An abstraction layer between storage formats and processing systems would reduce the $N \times M$ problem to $N + M$. Each system implements the interface once, and each format provides a

¹Apache Parquet, <https://parquet.apache.org> (accessed: May 22, 2026)

²Apache ORC, <https://orc.apache.org/> (accessed: May 22, 2026)

³PostgreSQL, <https://www.postgresql.org/> (accessed: May 22, 2026)

single implementation against it. Gienieczko et al. (2025) propose exactly this and present AnyBlox as their realization.

2 Self-Decoding Data: The AnyBlox Design

A future-proof data access solution must satisfy two conceptual requirements simultaneously. First, it must introduce an abstraction layer that decouples format internals from processing systems and system internals from decoders, thereby reducing the $N \times M$ integration burden to $N + M$. Second, this abstraction must not preclude direct file access, since routing data through an intermediary process introduces transfer overhead that can dominate total runtime. These two goals are in tension: an abstraction layer must be located somewhere, and every candidate location either breaks direct access or reintroduces format-specific coupling [3].

Existing approaches each resolve this tension in favor of one side. Native integration preserves direct file access and achieves the best performance through tight coupling with host internals, but requires a separate decoder implementation per format per system, thus reproducing the $N \times M$ problem. Extension mechanisms reduce the per-format implementation effort, but each host exposes its own extension API, so a decoder written for system A cannot be reused in system B without a near-complete rewrite. Isolated extensions, such as containerized user-defined functions [10] or remote cloud user-defined code execution solutions provide both isolation and portability, but introduce a data transfer bottleneck between the host and the isolated decoder process [3].

Given these constraints, Gienieczko et al. (2025) argue that there is only one location for the decoder that satisfies both requirements: inside the data file itself. Packaging the decoder with the data eliminates the transfer bottleneck, since decoder and data reside in the same address space, while the abstraction is preserved because the host need not implement any format-specific logic. This is the core idea behind self-decoding data and the design premise of AnyBlox.

2.1 WebAssembly as the Decoder Runtime

A self-decoding dataset requires a decoder runtime that satisfies four criteria: **Portability**, **Security**, **Performance** and **Extensibility**.

Gienieczko et al. (2025) identify WebAssembly (Wasm) as a runtime that satisfies all four simultaneously. Wasm is a portable low-level bytecode and virtual machine specification, designed as an abstraction over modern hardware rather than a commitment to any specific programming model, making it language-, hardware-, and platform-independent [4].

2.1.1 Portability. A decoder implemented in Wasm runs on a wide variety of host architectures (e.g. browsers, mobile devices, embedded, and server workstations) [3]. For AnyBlox, this means a format author writes and compiles a decoder once, and any host can execute it via `libanyblox`, independent of the underlying hardware or operating system.

2.1.2 Security. The WebAssembly specification provides machine-verifiable memory safety and isolation guarantees [12]. AnyBlox reinforces this by being implemented predominantly in safe Rust, guaranteeing memory and thread safety. All unsafe code is confined

to a single module handling Wasm memory maps, keeping it easily auditable [3].

2.1.3 Performance. Wasm is JIT-compiled to the host’s native instruction set at runtime, making near-native performance attainable. Additionally, AnyBlox avoids unnecessary data copying through a custom memory management scheme. Gienieczko et al. (2025) note that AnyBlox can passively benefit from future improvements in Wasm compiler technology, since all JIT and sandboxing details are encapsulated in `libanyblox` and not exposed in the public API [3].

2.1.4 Extensibility. Since most general-purpose programming languages provide a Wasm toolchain, format authors can implement decoders in their language of choice without restrictions. The main practical obstacle is SIMD instructions: converting platform-specific SIMD intrinsics to Wasm’s instruction set requires manual developer effort. However, since Wasm defines SIMD at the bytecode level, the resulting modules remain fully portable across architectures [3].

2.2 Reading a Dataset: From File to Query Pipeline

When a Host wants to read an AnyBlox dataset, it first opens the `.any` file and reads its metadata during query planning. The metadata provides the Host with the schema, the total number of rows, and a recommended batch size, all of which are needed to partition the workload across threads and schedule the scan before decoding begins. The Host then passes this metadata and the `.any` file to `libanyblox`, which extracts the Wasm Decoder module, JIT-compiles it, and caches the result (subsequent uses of the same Decoder therefore incur no compilation cost). AnyBlox then sets up the execution environment for the Decoder and returns a job handle to the Host.

From this point, the Host repeatedly requests arbitrary tuple ranges from that handle. Modern query engines process data in batches to amortize operator communication overhead, so rather than decoding the entire dataset at once, the Host issues successive calls for ranges of tuples [5]. AnyBlox delegates each such request into the sandboxed Wasm module and returns the decoded result to the Host. The sole function through which this happens is `decode_batch` [3]:

- `i32 data`: pointer to the start of the encoded data in Wasm linear memory
- `i32 data_length`: length of the encoded data in bytes
- `i32 start_tuple`: ID of the first tuple to decode
- `i32 tuple_count`: number of tuples to decode
- `i32 state`: pointer to the state page in Wasm linear memory
- `i64 projection_mask`: bitmask specifying which columns to decode, allowing the Host to skip columns that are not needed for the current query

The `data` parameter points to the encoded dataset inside the Decoder’s execution environment, but how does the dataset get there in the first place? A naive solution would be to copy the entire dataset into a memory region accessible to the Decoder, but this would incur a heavy initialization cost proportional to the dataset size. Instead, AnyBlox maps the dataset file directly into

the Decoder’s linear memory via `mmap`, a system call that maps a file directly into the virtual address space of a process without copying it, so physical memory is only allocated when a page is first accessed.

Wasm’s linear memory is a contiguous, byte-addressable memory region starting at address `0x0` that represents the entire address space accessible to the Decoder, isolated from the rest of the host process. AnyBlox controls its layout: the encoded dataset is mapped into it via `mmap`, and the state page is placed immediately after. The data parameter is therefore simply a pointer to the start of the mapped dataset within this region. AnyBlox reuses the linear memory region across successive jobs on the same thread, avoiding repeated initialization costs.

The state parameter points to a page within the linear memory that is zero-initialized at the start of every job. A Decoder can use it to cache metadata or memory allocations across successive `decode_batch` calls. For example, a run-end encoding Decoder that performs a binary search on the first call can store its current position on the state page and resume directly on the next call, skipping the search entirely. Because the state page resets at the start of each job, Decoders are idempotent: identical parameters on the same dataset always produce equivalent results.

Once decoding is complete, `decode_batch` returns a pointer to an Apache Arrow Record Batch within the linear memory. Arrow⁴ is a columnar in-memory format developed by the Apache Software Foundation: rather than storing data row by row, it organizes each column as a contiguous array in memory, which allows query engines to process large amounts of data efficiently without unnecessary memory accesses. A Record Batch is a collection of such arrays, all of the same length, representing a horizontal slice of a table. Arrow standardizes the representation of common SQL types (e.g., integers, floating-point numbers, strings, dates, and more) and is supported by a wide range of query engines. The choice of Arrow as the uniform output format of AnyBlox follows the same principle as other intermediate representations in systems design, because by agreeing on a common language between two components, the number of required translations reduces from $N \times M$ to $N + M$. Every Decoder, regardless of the underlying encoding scheme, produces an Arrow Record Batch, and every Host that speaks Arrow can consume it without any AnyBlox-specific logic. AnyBlox translates the Wasm-internal pointer into a native address and passes the batch to the Host, where it enters the query pipeline. The degree

of integration varies by system: DataFusion [7] uses Arrow as its primary internal representation throughout, making the transfer zero-copy end-to-end. DuckDB [8] and Spark⁵ do not use Arrow internally, but both have built-in support for reading Arrow batches and converting them into their respective internal representations on the fly. This conversion logic is part of the mainline code of both systems and requires no AnyBlox-specific adaptation [3].

References

- [1] Michael Armbrust, Ali Ghodsi, Reynold Xin, and Matei Zaharia. 2021. Lakehouse: A New Generation of Open Platforms That Unify Data Warehousing and Advanced Analytics. (2021).
- [2] E. F. Codd. 1970. A Relational Model of Data for Large Shared Data Banks. *Commun. ACM* 13, 6 (June 1970), 377–387. <https://doi.org/10.1145/362384.362685>
- ⁴Apache Arrow, <https://arrow.apache.org/> (accessed: June 2, 2026)
- ⁵Apache Spark, <https://spark.apache.org/> (accessed: June 2, 2026)
- [3] Mateusz Gieniec, Maxilian Kuschewski, Thomas Neumann, Viktor Leis, and Jana Giceva. 2025. AnyBlox: A Framework for Self-Decoding Datasets. *Proceedings of the VLDB Endowment* 18, 11 (July 2025), 4017–4031. <https://doi.org/10.14778/3749646.3749672>
- [4] Andreas Haas, Andreas Rossberg, Derek L. Schuff, Ben L. Titzer, Michael Holman, Dan Gohman, Luke Wagner, Alon Zakai, and JF Bastien. 2017. Bringing the Web up to Speed with WebAssembly. *SIGPLAN Not.* 52, 6 (June 2017), 185–200. <https://doi.org/10.1145/3140587.3062363>
- [5] Timo Kersten, Viktor Leis, Alfons Kemper, Thomas Neumann, Andrew Pavlo, and Peter Boncz. 2018. Everything You Always Wanted to Know about Compiled and Vectorized Queries but Were Afraid to Ask. *Proceedings of the VLDB Endowment* 11, 13 (Sept. 2018), 2209–2222. <https://doi.org/10.14778/3275366.3275370>
- [6] Laurens Kuiper. 2025. Query Engines: Gatekeepers of the Parquet File Format. <https://duckdb.org/2025/01/22/parquet-encodings.html>.
- [7] Andrew Lamb, Yijie Shen, Daniel Heres, Jayjeet Chakraborty, Mehmet Ozan Kabak, Liang-Chi Hsieh, and Chao Sun. 2024. Apache Arrow DataFusion: A Fast, Embeddable, Modular Analytic Query Engine. In *Companion of the 2024 International Conference on Management of Data (SIGMOD ’24)*. Association for Computing Machinery, New York, NY, USA, 5–17. <https://doi.org/10.1145/3626246.3653368>
- [8] Mark Raasveldt and Hannes Mühleisen. 2019. DuckDB: An Embeddable Analytical Database. In *Proceedings of the 2019 International Conference on Management of Data (SIGMOD ’19)*. Association for Computing Machinery, New York, NY, USA, 1981–1984. <https://doi.org/10.1145/3299869.3320212>
- [9] Red Hat, Inc. 2017. CVE-2017-15098. <https://nvd.nist.gov/vuln/detail/CVE-2017-15098>.
- [10] Karla Saur, Tara Mirmira, Konstantinos Karanasos, and Jesús Camacho-Rodríguez. 2022. Containerized Execution of UDFs: An Experimental Evaluation. *Proceedings of the VLDB Endowment* 15, 11 (July 2022), 3158–3171. <https://doi.org/10.14778/3551793.3551860>
- [11] Abraham Silberschatz, Henry F. Korth, and S. Sudarshan. 2020. *Database System Concepts* (seventh edition ed.). McGraw-Hill, New York, NY.
- [12] Conrad Watt. 2018. Mechanising and Verifying the WebAssembly Specification. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2018)*. Association for Computing Machinery, New York, NY, USA, 53–65. <https://doi.org/10.1145/3167082>